

# ICP4

Student Name: Srinivas Musinuri

Student id:700758813

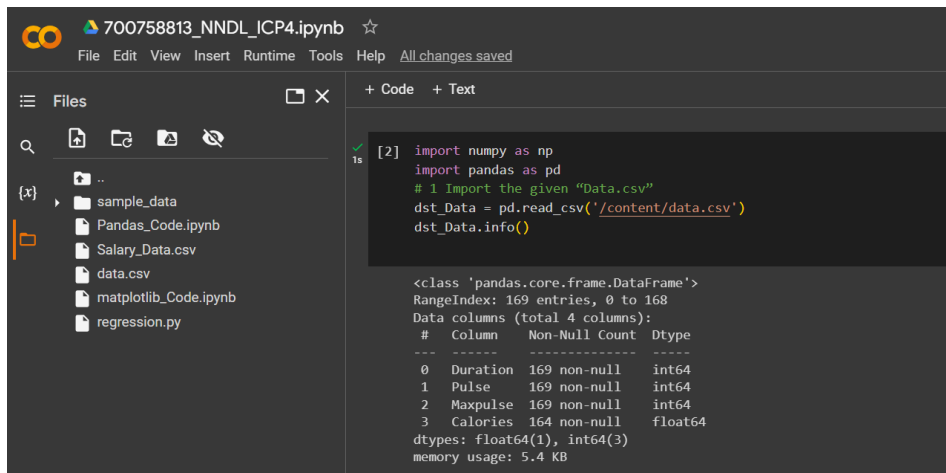
GitHub Link: [https://github.com/srinivasmusinuri/700758813\\_NNDL\\_ICP4](https://github.com/srinivasmusinuri/700758813_NNDL_ICP4)

Video Link:

[https://drive.google.com/file/d/1tdcyMJxzvThn85YoVXFr4buGFZMn\\_8sj/view?usp=sharing](https://drive.google.com/file/d/1tdcyMJxzvThn85YoVXFr4buGFZMn_8sj/view?usp=sharing)

## 1. Data Manipulation

- Read the provided CSV file 'data.csv'.
- <https://drive.google.com/drive/folders/1h8C3mLsso-R-siOLsvoYwPLzy2fJ4IOF?usp=sharing>



The screenshot shows a Jupyter Notebook interface with a file explorer on the left and a code editor on the right. The file explorer shows a directory named 'sample\_data' containing files like 'Pandas\_Code.ipynb', 'Salary\_Data.csv', 'data.csv', 'matplotlib\_Code.ipynb', and 'regression.py'. The code editor shows the following code:

```
[2]: import numpy as np
import pandas as pd
# 1 Import the given "Data.csv"
dst_Data = pd.read_csv('/content/data.csv')
dst_Data.info()
```

The output of the code is displayed below the code cell:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 169 entries, 0 to 168
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   Duration    169 non-null    int64
1   Pulse       169 non-null    int64
2   Maxpulse    169 non-null    int64
3   Calories    164 non-null    float64
dtypes: float64(1), int64(3)
memory usage: 5.4 KB
```

- Show the basic statistical description about the data.
- Check if the data has null values.
  - Replace the null values with the mean.

700758813\_NNDL\_ICP4.ipynb ☆

File Edit View Insert Runtime Tools Help All changes saved

Files

- sample\_data
  - Pandas\_Code.ipynb
  - Salary\_Data.csv
  - data.csv
  - matplotlib\_Code.ipynb
  - regression.py

+ Code + Text

```
# Show the basic statistical description about the data.
dst_Data.head()
```

	Duration	Pulse	Maxpulse	Calories
0	60	110	130	409.1
1	60	117	145	479.0
2	60	103	135	340.0
3	45	109	175	282.4
4	45	117	148	406.0

```
[4] # Check if the data has null values.
dst_Data.isnull().any()
```

```
Duration    False
Pulse       False
Maxpulse    False
Calories    True
dtype: bool
```

700758813\_NNDL\_ICP4.ipynb ☆

File Edit View Insert Runtime Tools Help All changes saved

Files

- sample\_data
  - Pandas\_Code.ipynb
  - Salary\_Data.csv
  - data.csv
  - matplotlib\_Code.ipynb
  - regression.py

+ Code + Text

```
[5] dst_Data.fillna(dst_Data.mean(), inplace=True)
dst_Data.isnull().any()
```

```
Duration    False
Pulse       False
Maxpulse    False
Calories    False
dtype: bool
```

```
[6] # Replace the null values with the mean
column_means = dst_Data.mean()
print(column_means)
dst_Data = dst_Data.fillna(column_means)
print(dst_Data.head(20))
```

```
Duration    63.846154
Pulse       107.461538
Maxpulse    134.047337
Calories    375.790244
dtype: float64
```

	Duration	Pulse	Maxpulse	Calories
0	60	110	130	409.100000
1	60	117	145	479.000000
2	60	103	135	340.000000
3	45	109	175	282.400000
4	45	117	148	406.000000

Disk 81.43 GB available

e. Select at least two columns and aggregate the data using: min, max, count, mean.

700758813\_NNDL\_ICP4.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Files

- sample\_data
- Pandas\_Code.ipynb
- Salary\_Data.csv
- data.csv
- matplotlib\_Code.ipynb
- regression.py

+ Code + Text

```
[6] 13 60 104 132 379.300000
14 60 98 123 275.000000
15 60 98 120 215.200000
16 60 100 120 300.000000
17 45 90 112 375.790244
18 60 103 123 323.000000
19 45 97 125 243.000000
```

```
[7] # Select at least two columns and aggregate the data using: min, max, count, mean.
res = dst_Data.agg({'Calories': ['mean', 'min', 'max', 'count'], 'Pulse': ['mean', 'min', 'max', 'count']})
print(res)
```

	Calories	Pulse
mean	375.790244	107.461538
min	50.300000	80.000000
max	1860.400000	159.000000
count	169.000000	169.000000

f. Filter the dataframe to select the rows with calories values between 500 and 1000.

g. Filter the dataframe to select the rows with calories values > 500 and pulse.

700758813\_NNDL\_ICP4.ipynb

File Edit View Insert Runtime Tools Help All changes saved

Files

- sample\_data
- Pandas\_Code.ipynb
- Salary\_Data.csv
- data.csv
- matplotlib\_Code.ipynb
- regression.py


+ Code + Text

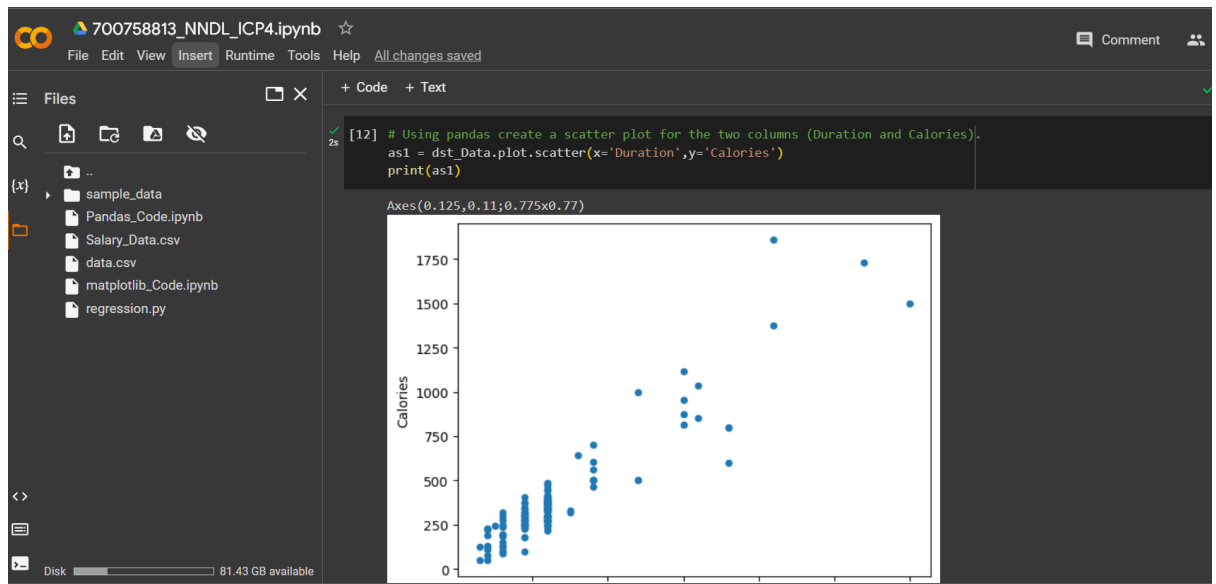
```
[8] # Filter the dataframe to select the rows with calories values between 500 and 1000
filter_dst_Data1=dst_Data[(dst_Data['Calories'] > 500) & (dst_Data['Calories'] < 1000)]
print(filter_dst_Data1)
# Filter the dataframe to select the rows with calories values > 500 and pulse < 100.
filter_dst_Data2=dst_Data[(dst_Data['Calories'] > 500) & (dst_Data['Pulse'] < 100)]
print(filter_dst_Data2)
```

	Duration	Pulse	Maxpulse	Calories
51	80	123	146	643.1
62	160	109	135	853.0
65	180	90	130	800.4
66	150	105	135	873.4
67	150	107	130	816.0
72	90	100	127	700.0
73	150	97	127	953.2
75	90	98	125	563.2
78	120	100	130	500.4
90	180	101	127	600.1
99	90	93	124	604.1
103	90	90	100	500.4
106	180	90	120	800.3
108	90	90	120	500.3
65	180	90	130	800.4
70	150	97	120	1115.0
73	150	97	127	953.2
75	90	98	125	563.2
65	60	93	122	604.1

k. Using pandas create a scatter plot for the two columns (Duration and Calories).

Example





## 2. Linear Regression

- Import the given "Salary\_Data.csv"
- Split the data in train\_test partitions, such that 1/3 of the data is reserved as test subset.
- Train and predict the model.
- Calculate the mean\_squared error

```
700758813_NNDL_ICP4.ipynb
File Edit View Insert Runtime Tools Help All changes saved

Files
sample_data
Pandas_Code.ipynb
Salary_Data.csv
data.csv
matplotlib_Code.ipynb
regression.py

Code
[12] # 2. Import the given "Salary_Data.csv"
dst_Sal = pd.read_csv('/content/Salary_Data.csv')
dst_Sal.info()
dst_Sal.head()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30 entries, 0 to 29
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   YearsExperience  30 non-null    float64
1   Salary          30 non-null    float64
dtypes: float64(2)
memory usage: 608.0 bytes

YearsExperience  Salary
0               1.1  39343.0
1               1.3  46205.0
2               1.5  37731.0
3               2.0  43525.0
4               2.2  39891.0

[14] A = dst_Sal.iloc[:, :-1].values #excluding last column i.e., years of experience column
B = dst_Sal.iloc[:, 1].values #only salary column

[15] # Split the data in train_test partitions, such that 1/3 of the data is reserved as test subset.
from sklearn.model_selection import train_test_split
A_train, A_test, B_train, B_test = train_test_split(A, B, test_size=1/3, random_state=0)

[16] # Train and predict the model.
from sklearn.linear_model import LinearRegression
reg = LinearRegression()
reg.fit(A_train, B_train)
B_Pred = reg.predict(A_test)

array([ 48835.18590871, 123079.39948819,  65134.55626083,  63265.36777221,
        115682.64545369, 108125.8914992 , 116537.23969801,  64199.96281652,
         76349.68719258, 100649.1375447 ])

[17] # Calculate the mean_squared error
S_error = (B_Pred - B_test) ** 2
Sum_Serror = np.sum(S_error)
mean_squared_error = Sum_Serror / B_test.size
mean_squared_error

21026037.329511296
```

- Visualize both train and test data using scatter plot.

